



# SQL

## Data Definition Language

André Restivo

# Index

Introduction

Table Basics

Data Types

Defaults

Constraints

Check

Not Null

Primary Keys

Unique Keys

Foreign Keys

Sequences

# Introduction

# SQL

- Structured Query Language.
- A special purpose language to manage data stored in a **relational** database.
- Based on **relational algebra**.
- Pronounced *Sequel*

# History

- Early 70's SEQUEL Developed at IBM
- 1986 SQL-86 and SQL-87 Ratified by ANSI and ISO.
- 1989 SQL-89
- 1992 SQL-92 Also know as SQL2.
- 1999 SQL:1999 Also known as SQL3 Includes regular expressions, recursive queries, triggers, non-scalar data types and some object-oriented expressions.
- 2003 SQL:2003 XML support and auto-generated values.
- 2006 SQL:2006 XQuery support.
- 2008 SQL:2008.
- 2011 SQL:2011.

# Standard

- Although SQL is an ANSI/ISO standard, every database system implements it in a slightly different way.
- These slides will try to adhere to the standard as much as possible.
- Sometimes we'll deviate and talk specifically about PostgreSQL.

# Table Basics

# Creating Tables

The basic structure of a table creation statement in SQL:

```
CREATE TABLE <table_name> (  
    <column_name> <data_type>,  
    <column_name> <data_type>,  
    ...  
    <column_name> <data_type>  
);
```

? Values between <> are to be replaced.



# Deleting Tables

To delete a table we do:

```
DROP TABLE <table_name>;
```

If there are foreign keys referencing the table we must use the *cascade* keyword:

```
DROP TABLE <table_name> CASCADE;
```

# Data Types

## PostgreSQL

# Numeric

- The types `SMALLINT`, `INTEGER`, and `BIGINT` store whole numbers.
- `NUMERIC(precision, scale)` stores numbers with a very large number of digits. The *scale* of a numeric data type is the count of decimal digits in the fractional part, to the right of the decimal point. The *precision* of a numeric is the total count of significant digits in the whole number.
- The data types `REAL` and `DOUBLE` precision are inexact, variable-precision numeric types.

# Numeric Precision

Precision for numeric types:

SMALLINT	2 bytes	-32768 to +32767
INTEGER	4 bytes	-2147483648 to +2147483647
BIGINT	8 bytes	-9223372036854775808 to +9223372036854775807
NUMERIC	user-specified	numeric(3,2) -99.9 to 99.9
REAL	4 bytes	inexact 6 decimal digits precision
DOUBLE	8 bytes	inexact 15 decimal digits precision

# Date / Time

- The type **DATE** stores date values in the ISO 8601 format.

? Example: 1980-12-25.

- The type **TIME** stores time values in the ISO 8601 format.

? Examples: 04:05 or 04:05:06.789

- The type **TIMESTAMP** store time and date values in the ISO 8601 format.

? Example: 1980-12-25 04:05:06.789.

- Date/time types can also store timezone information.

# Text

- The type `CHARACTER VARYING(n)`, or `VARCHAR(n)`, stores variable-length text with a user defined limit.
- The type `CHARACTER(n)`, or `CHAR(n)`, stores fixed-length, blank-padded text.
- The type `TEXT` stores variable unlimited length text.

? `VARCHAR`, without a limit, is the same as `TEXT`

? `CHAR`, without a limit, is the same as `CHAR(1)`

# Boolean

- The BOOLEAN type can only have two possible values: `true` or `false`.
- Possible values representing true: `TRUE`, `'t'`, `'true'`, `'y'`, `'yes'`, `'on'` and `'1'`
- Possible values representing false: `FALSE`, `'f'`, `'false'`, `'n'`, `'no'`, `'off'` and `'0'`

# Serial

- In PostgreSQL there is a *pseudo-type* that can be used to define **auto-generated** identifiers or **auto-counters**.
- To define a column as an auto-counter we use the type **SERIAL**.



# Example

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128),  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2),  
  taxes NUMERIC(6,2),  
  card_number INTEGER,  
  active BOOLEAN  
);
```

# Defaults

# Default Values

For each column we can define its default value:

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> DEFAULT <default_value>,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>  
);
```

? The default default value is NULL

# Example

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128),  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0,  
  active BOOLEAN DEFAULT TRUE  
);
```

? Card number default value is 0.

? Employee is active by default.

# Constraints

# Constraint Types

Several types of constraints can be defined using SQL:

- Check
- Not Null
- Unique Keys
- Primary Keys
- Foreign Key

Constraints can be **column based** or **table based**.

- **Column** constraints appear in front of the column they are referring to.
- **Table** constraints appear as a **separate** clause.

Check

# Check Constraint

Check constraints allows to define a constraint on the column values using an expression that the values must follow:

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> CHECK <check_expression>,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>  
);
```



# Example

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128),  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2) CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE  
);
```

? Salary must be larger than 500.

? Card number must be larger or equal to 0.

# Constraint Naming

Giving names to constraints allows us to better identify them when errors occur or when we want to refer to them.

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> CONSTRAINT <constraint_name> CHECK <check_expression>,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>  
);
```

? Naming constraints is optional but is a good practice.

# Example

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128),  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE  
);
```

# Multiple Column Check

If the check constraint refers to more than one column, we must use a table based constraint:

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128),  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? Taxes have to be lower than salary.

**Not Null**

# Not Null Constraint

We can define that a certain column does not allow NULL values:

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> NOT NULL,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>  
);
```

# Example

```
CREATE TABLE employee (  
  id SERIAL,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? Name cannot be null.

# Primary Keys



# Primary Key Constraints

We can define one, and only one, primary key for our table:

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> PRIMARY KEY,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>  
);
```

# Example

```
CREATE TABLE employee (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? Id cannot be null and cannot have repeated values.

# Multiple Column Primary Key

If a primary key is composed by more than one column, we must use a **table based** constraint:

```
CREATE TABLE <table_name> (  
    <column_name> <data_type>,  
    <column_name> <data_type>,  
    ...  
    <column_name> <data_type>,  
    PRIMARY KEY (<column_name>, <column_name>)  
);
```

# Example

```
CREATE TABLE telephone (  
  employee INTEGER,  
  phone VARCHAR,  
  PRIMARY KEY (employee, phone)  
);
```

? An employee cannot have the same phone number twice.

# Unique Keys

# Unique Key Constraints

Unique keys are identical to primary keys but:

- they **allow** NULL values;
- and there can be **multiple unique keys** in one table.

They are created using the same type of syntax:

# Example

```
CREATE TABLE employee (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER  
    UNIQUE  
    DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? Card keys cannot have repeated values.

# Multiple Column Unique Keys

If a unique key is composed by more than one column, we must use a **table based** constraint:

```
CREATE TABLE <table_name> (  
  <column_name> <data_type>,  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>,  
  UNIQUE (<column_name>, <column_name>)  
);
```



# Foreign Keys

# Foreign Key Constraints

- We can also declare foreign keys.
- A foreign key must always **reference a key** (primary or unique) from another (or the same) table.
- Databases don't allow columns with a foreign key containing values that do not exist in the referenced column.

```
CREATE TABLE <table_A> (  
  <column_A> <data_type> PRIMARY KEY,  
  <column_B> <data_type>,  
  ...  
  <column_C> <data_type>  
);  
  
CREATE TABLE <table_B> (  
  <column_X> <data_type> PRIMARY KEY,  
  <column_Y> <data_type>,  
  ...  
  <column_Z> <data_type> REFERENCES <table_A>(<column_A>)  
);
```

# Example

```
CREATE TABLE employee (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER  
    UNIQUE  
    DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  department_id integer REFERENCES department(id),  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? The id of the department references the id column in the department table.

? Employees cannot have a department number that doesn't exist in the department table.

# Foreign Key to Primary Key

If the referenced column is the primary key of the other table, we can omit the name of the column:

```
CREATE TABLE employee (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER  
    UNIQUE  
    DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  department_id integer REFERENCES department,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? The id of the department references the primary key in the department table.

# Multiple Column Foreign Key

If the referenced table has a key with multiple columns, we must use a **table based** constraint to define our foreign key:

```
CREATE TABLE telephone (  
    employee INTEGER,  
    phone VARCHAR,  
    PRIMARY KEY (employee, phone)  
);  
  
CREATE TABLE call (  
    id INTEGER PRIMARY KEY,  
    employee INTEGER,  
    phone INTEGER,  
    when DATE,  
    caller INTEGER,  
    FOREIGN KEY (employee, phone) REFERENCES telephone (employee, phone)  
);
```

# Example

We can also omit the referenced columns if they are the primary keys:

```
CREATE TABLE telephone (  
    employee INTEGER,  
    phone VARCHAR,  
    PRIMARY KEY (employee, phone)  
);  
  
CREATE TABLE call (  
    id INTEGER PRIMARY KEY,  
    employee INTEGER,  
    phone INTEGER,  
    when DATE,  
    caller INTEGER,  
    FOREIGN KEY (employee, phone) REFERENCES telephone  
);
```

# Deleting Referenced Values

- Declaring a foreign key means that values in one table must also appear in the referenced column.
- When a line having values referencing it is deleted, three different things can occur:
  - An **error** is thrown.
  - The referencing values becomes **NULL**.
  - All referencing **lines are deleted** (this might cause a cascade effect).

# Updating Referenced Values

- The same problem happens when we update a line that is referenced by another column.
- When a line having values referencing it is updated, three different things can occur:
  - An **error** is thrown.
  - The referencing values becomes **NULL**.
  - All referencing **values are updated** to the new value (this might cause a cascade effect).



# On Delete and On Update

To define the desired behavior, we should use the `ON DELETE` and `ON UPDATE` clauses with one of three possible values:

- `RESTRICT` (throws an error)
- `SET NULL` (values become null)
- `CASCADE` (lines are deletes or values updated)

? `RESTRICT` is the default value.

# On Delete and On Update

```
CREATE TABLE <table_A> (  
  <column_A> <data_type> PRIMARY KEY,  
  <column_B> <data_type>,  
  ...  
  <column_C> <data_type>  
);  
  
CREATE TABLE <table_B> (  
  <column_X> <data_type> PRIMARY KEY,  
  <column_Y> <data_type>,  
  ...  
  <column_Z> <data_type> REFERENCES <table_A>(<column_A>)  
    ON DELETE SET NULL ON UPDATE CASCADE  
);
```

# Example

```
CREATE TABLE employee (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(128) NOT NULL,  
  address VARCHAR(256),  
  birthdate DATE,  
  salary NUMERIC(6,2)  
    CONSTRAINT minimum_wage CHECK (salary > 500),  
  taxes NUMERIC(6,2),  
  card_number INTEGER  
    UNIQUE  
    DEFAULT 0 CHECK (card_number >= 0),  
  active BOOLEAN DEFAULT TRUE,  
  department_id integer REFERENCES department  
    ON DELETE SET NULL ON UPDATE CASCADE,  
  CONSTRAINT taxes_lower_salary CHECK (taxes < salary)  
);
```

? If a department with id 1 is deleted, all employees with department id equal to 1 will start having a null department number.

? If a department with id 1 is updated to id 2, all employees with department id equal to 1 will start having a department number equal to 2.

# Sequences

# Sequences

- A sequence is a special kind of database object designed for generating unique numeric identifiers.
- They ensure that a different value is generated for every client.

# Sequences

```
CREATE SEQUENCE <name>;
```

Get the next and the current value of the sequence.

```
SELECT nextval('<name>');  
SELECT currval('<name>');
```

Getting the current value only works if called after calling nextval and in the same transaction.

# The SERIAL type

The data type serial is not a true type, but merely a notational convenience for setting up unique identifier columns.

```
CREATE TABLE <tablename> (  
    <colname> SERIAL  
);
```

Is equivalent to:

```
CREATE SEQUENCE <tablename_colname_seq>;  
CREATE TABLE <tablename> (  
    <colname> integer DEFAULT nextval('<tablename_colname_seq>') NOT NULL  
);
```

# Example

```
CREATE TABLE category (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR  
)  
  
CREATE TABLE product (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR,  
  cat_id INTEGER REFERENCES category  
);  
  
INSERT INTO category VALUES(DEFAULT, 'Fruits');  
INSERT INTO products VALUES(DEFAULT, 'Lemon', currval('category_id_seq'));
```